

# HW 4: CNN Image Categorization

David E. Farache, Email ID: dfarache@purdue.edu

March 7, 2023

## 1 Introduction

The goal of this homework is to create a pizza identifier with the implement of skip blocks within the network into a convolutional neural network (CNN). Skip-block block functions by adding the prior importance to the end of a network during backward propagation in order to increase the effect of the loss calculated. The evaluation of this networks must be done using MSE loss regression and Complete Box IOU loss.

## 2 Methodology

COCO was utilized to create an initial training set via the `pycocotools.coco` function, which would reduce to 3790 for training and 1987 for validation. The data-loader was based on HW4 with the addition of inter encoding based on the category and transform the image from a PIL image to a normalized tensor.

Across the course we have utilized CNN for predicting properties and are now expanding to ever deeper networks but these CNNs have an issue that with increasing depth there is degradation or a vanishing gradient. Skip-blocks resolve this issue by passing information from the initial layers to those deeper in the network, henceforth creating a "skip" or "short-cut". In this case our skip-block is a really a resnet block as the information is passed via matrix addition.

The CNN model was based on the DLStudio network with reduction of linear layer due to gpu space limitations. For this BatchNorm2d, the Skip-Block that was created, and Con2d for convolutional layers. The network had for loops that created layers of skip of depth 8.

For loss calculation, the classification is measured using Cross-Entropy loss while the bounding box is based upon the MSE regression loss.

$$Cross - Entropy = - \sum_{c=1}^M y_{o,c} \log(p_{o,c}) \quad (1)$$

$$MSE = \sum_{i=1}^D (x_i - y_i)^2 \quad (2)$$

For boundary box analysis, IoU is typically utilized which measure loss by taking 1 minus the area of overlap by the area of union. GIoU came about in order to grant an error even when there is no intersection so that the model could move into the proper direction and not keep guessing until it overlaps with the box. The issue was that although GIoU helped find the box, it had issue fitting.

Complete Box IoU was mode to combat this issue by introducing aspect ratio between bonding boxes so that the shape matches more so that being the right position.

$$CIoULoss = 1 - IoU + \frac{|C - B \cup B^{gt}|}{|C|} + \frac{d^2}{c^2} + \alpha * v \quad (3)$$

### 3 Task 1: COCO Data-set Preparation

For task 1, the COCO data-set was used from train2014 and val2014 in order to reduce and create a training and validation for categories: pizza, cat, and bus. Images were selected based on whether there was a single dominant label with area greater than 200 x 200. Each image was turned in RGB and resized from their original height to 256 x 256. The bbox of the image was also re-scaled. A data-frame of filename, path, category, and bbox dimensions was saved for efficient processing.

```

1 # %% [markdown]
2 # # Library
3
4 # %%
5 %matplotlib inline
6 from pycocotools.coco import COCO
7
8 import numpy as np
9 import skimage.io as io
10 import skimage
11 import cv2
12 import pandas as pd
13
14 import matplotlib.pyplot as plt
15 import pylab
16 import random
17 from PIL import Image
18 pylab.rcParams['figure.figsize'] = (8.0, 10.0)
19
20 # %% [markdown]
21 # # Display Images
22
23 # %%
24 # Based from skeleton code given
25 def display_random_image_with_bbox(new_image_path, filename, bbox, cat):
26     file = new_image_path + '/' + filename
27     image = Image.open(file)
28     [x, y, w, h] = bbox
29     fig, ax = plt.subplots(1,1)
30     image = np.uint8(image)
31     image = cv2.rectangle(image, (int(x), int(y)), (int(x + w), int(y + h)),
32 color=(36, 255, 12), thickness=2)
33     image = cv2.putText(image, cat, (int(x), int(y - 10)), fontFace=cv2.
34 FONT_HERSHEY_SIMPLEX, fontScale=0.8, color=(36, 255, 12), thickness=2)
35
36     ax.imshow(image)
37     ax.set_axis_off()
38     plt.axis("tight")
39     plt.show()
40
41     fig, axs = plt.subplots(nrows=3, ncols=3, figsize=(10,10), dpi = 150)

```

```

40     axs = axs.flatten()
41     axs_count = 0
42
43 # %%
44 def display_images(saveForPlotting):
45     fig, ax = plt.subplots(3, 3)
46     row, col = 0, 0
47
48     #print(len(saveForPlotting))
49     for arr in saveForPlotting:
50         file = arr[0] + '/' + arr[1]
51         image = Image.open(file)
52
53         [x, y, w, h] = arr[2]
54         image = np.uint8(image)
55         image = cv2.rectangle(image, (int(x), int(y)), (int(x + w), int(y + h)),
56                               color=(36, 255, 12), thickness=2)
57         image = cv2.putText(image, arr[3], (int(x), int(y - 10)), fontFace=cv2.
58                               FONT_HERSHEY_SIMPLEX, fontScale=0.8, color=(36, 255, 12), thickness=2)
59         ax[row, col].imshow(image)
60
61         # Increment through row
62         col += 1
63
64         if col == 3:
65             col = 0
66             row += 1
67
68     fig, axs = plt.subplots(nrows=3, ncols=3, figsize=(10,10), dpi = 150)
69     plt.show()
70     axs = axs.flatten()
71
72 # %% [markdown]
73 # # Adjust and Change Images
74
75 # %%
76 def resizeBBOX(bbox, startWidth, startHeight, resize):
77     x_scale = resize / startWidth
78     y_scale = resize / startHeight
79     bboxResize = np.zeros(4)
80
81     bboxResize[0] = int(x_scale * bbox[0])
82     bboxResize[1] = int(y_scale * bbox[1])
83     bboxResize[2] = int(x_scale * bbox[2])
84     bboxResize[3] = int(y_scale * bbox[3])
85
86     return bboxResize.tolist()
87
88 # %%
89 def resizeAndRepairImage(start_image_path, new_image_path, filename, bbox, width
90                           , height, resize):
91     img = Image.open(start_image_path + '/' + filename)
92     width, height = img.size
93
94     image = img.convert(mode="RGB")
95
96     img = image.resize((resize, resize), Image.BOX)

```

```

94         img.save(new_image_path + '/' + filename)
95
96 # %%
97 def saveAsDataFrame(ids, cats, filepaths, x1, y1, width, height, type):
98     # Containers for data to go in training label csv
99     columns = ["id", "category", "filepath", "x1", "y1", "width", "height"]
100    dataframe = pd.DataFrame(columns=columns)
101    dataframe["id"] = ids
102    dataframe["category"] = cats
103    dataframe["filepath"] = filepaths
104    dataframe["x1"] = x1
105    dataframe["y1"] = y1
106    dataframe["width"] = width
107    dataframe["height"] = height
108    dataframe.to_csv("{0}_labels.csv".format(type))
109
110
111 # %% [markdown]
112 # # Choose Images Within Parameters
113
114 # %%
115 # Get Random Images from set
116 def ImageSelection(start_image_path, new_image_path, cocoObj, class_list, type,
117                    spotSave):
118     # Save File Location List
119     saveImportant = []
120     saveForPlotting = []
121
122     # Save image info
123     ids = []
124     cats = []
125     filepaths = []
126     x1 = []
127     y1 = []
128     widthSave = []
129     heightSave = []
130
131     for cat in class_list:
132         # get all images containing given categories
133         catIds = cocoObj.getCatIds(catNms=[cat]) # Get ids from annotations
134         imgIds = cocoObj.getImgIds(catIds=catIds) # Load images ids of chosen
135         annotations ids
136         img = cocoObj.loadImgs(ids=imgIds) # Get images
137         numPlots = 0 # Make X number plots
138
139         #Loop per image
140         for idx, images in enumerate(img):
141             annIds = cocoObj.getAnnIds(imgIds=images['id'], catIds=catIds,
142             iscrowd=False) # Get dictionary value
143             anns = cocoObj.loadAnns(annIds) # Get annotations
144             domObj = 0 # Amount of Dominante Obj
145             spot = 0
146             for jdx, ann in enumerate(anns):
147                 if ann['area'] > 200 * 200: #check if means parameter of
148                     dominate obj
149                     domObj += 1

```

```

147         spot = jdx
148
149
150         if domObj == 1: # only one dom allowed
151             # Resize Images
152             width, height = images['width'], images['height']
153             boxDict = []
154             #for ann in anns:
155             bboxResize = resizeBBOX(anns[jdx]['bbox'], int(width), int(
height), resize=256) #Adjust Box
156             resizeAndRepairImage(start_image_path, new_image_path, images['
file_name'], anns[jdx]['bbox'], width, height, resize=256) # Save Adjust
Image
157
158             ids.append(images['id'])
159             cats.append(cat)
160             filepaths.append("{0}/{1}".format(spotSave, images['file_name']))
161             x1.append(bboxResize[0])
162             y1.append(bboxResize[1])
163             widthSave.append(bboxResize[2])
164             heightSave.append(bboxResize[3])
165
166             if numPlots < 3:
167                 #display_random_image_with_bbox(new_image_path, images['
file_name'], bboxResize, cat)
168                 saveForPlotting.append([new_image_path, images['file_name'],
bboxResize, cat])
169                 print(saveForPlotting)
170                 numPlots += 1
171
172             display_images(saveForPlotting)
173             saveAsDataFrame(ids, cats, filepaths, x1, y1, widthSave, heightSave, type)
174
175 # %%
176 # Input
177 train_json = '/Users/davidfarache/Documents/ECE60146/HW5/annotations/
instances_train2014.json'
178 val_json = '/Users/davidfarache/Documents/ECE60146/HW5/annotations/
instances_val2014.json'
179
180 train_path = '/Users/davidfarache/Documents/ECE60146/HW5/train2014'
181 train_data_path = '/Users/davidfarache/Documents/ECE60146/HW5/trainingData'
182
183 val_path = '/Users/davidfarache/Documents/ECE60146/HW5/val2014'
184 val_data_path = '/Users/davidfarache/Documents/ECE60146/HW5/valData'
185
186 trainSaveSpot = 'trainingData'
187 valSaveSpot = 'valData'
188
189 class_list = ['pizza', 'bus', 'cat']
190
191 # %%
192 cocoTrain = COCO(train_json)
193 cocoVal = COCO(val_json)
194
195 ImageSelection(train_path, train_data_path, cocoTrain, class_list, 'train',
trainSaveSpot)

```

```

196 ImageSelection(val_path, val_data_path, cocoVal, class_list, 'val', valSaveSpot)
197
198 # %%

```

**Listing 1:** *Creating COCO data-set*

A sample of the images select has been placed below, which can be seen to have shrunk and had their boundary box rescaled within them:



**Figure 1:** *Selection of Training Set*

## 4 Task 2: Create Skip-Connection Block and CNN

The network created utilizes a training and validation set created from COCO. The following code has the data-set which stores boundary box, normalized image, and the image label. The depth was set to 8, with learning rate of 1e-4, trained for 7 epochs, and running average loss value is taken per 100 iterations. Two training functions were created one that performed Cross-Entropy and MSE loss (trainMSERegression) and the other using Cross-Entropy and CIoU (trainCompleteBoxIOU). The Adam optimizer was selected for both. The testing was done the same for both.

```

1 # %%
2 # Import Libraries
3 import numpy as np
4
5 # PyTorch
6 import torch
7 import torchvision.transforms as tvt
8 import torch.utils.data
9 import torch.nn as nn
10 import torch.nn.functional as F
11 from torchvision.ops import complete_box_iou_loss
12
13 # Data Processing
14 from PIL import Image
15 import os
16 import cv2
17 import pandas as pd
18
19 # Plotting
20 import matplotlib.pyplot as plt
21 import seaborn as sns
22
23 # GLOBAL VARIABLES
24 device = 'cuda' if torch.cuda.is_available() else 'cpu'
25 device = torch.device(device)
26
27 # %% [markdown]
28 # # Save Annotations into Dictionary
29
30 # %%
31 def ImageProcessing(images):
32     # Get dir
33     image_dir = images["filepath"]
34
35     # Get inputs
36     image = Image.open(image_dir)
37     bbox = [images["x1"], images["y1"], images["width"], images["height"]]
38
39     # Normalize image
40     toTensor = tvt.ToTensor()(image)
41     toNormalize = tvt.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))(toTensor)
42     return toNormalize, bbox
43
44 # %%
45 # Create Dataset
46 class MyDataset(torch.utils.data.Dataset):
47     # Obtain meta information (e.g. list of file names)
48     def __init__(self, imagesDataFrame, class_list):
49         super().__init__()
50         self.imagesDataFrame = imagesDataFrame
51         self.catEncoding = { # Set Integer Values for Cat
52             "bus" : 0,
53             "cat" : 1,
54             "pizza" : 2
55         }
56
57     def __len__(self):

```

```

58         return len(self.imagesDataFrame)
59
60     def __getitem__(self, i):
61         imagesDataFrame = self.imagesDataFrame.iloc[i]
62         normalImage, bbox = ImageProcessing(imagesDataFrame)
63         label = self.catEncoding[imagesDataFrame["category"]]
64
65         # Fix box of resize
66         width, height = bbox[2], bbox[3]
67         bboxAdjust = [bbox[0], bbox[1], width, height]
68         bboxTensor = torch.tensor(bboxAdjust, dtype=torch.float)
69         return normalImage, label, bboxTensor
70
71 # %% [markdown]
72 # # Skip Block
73
74 # %%
75 # Based on DLStudio SkipBlock Code
76 # Based off following https://blog.paperspace.com/writing-resnet-from-scratch-in-pytorch/
77 class Block(nn.Module):
78     def __init__(self, in_ch, out_ch, downsample=False, skip_connections=True)
79     :
80         super(Block, self).__init__()
81         self.downsample = downsample
82         self.in_ch = in_ch
83         self.out_ch = out_ch
84
85         self.skip_connections = skip_connections
86
87         self.conv1 = nn.Sequential(
88             nn.Conv2d(in_ch, out_ch, kernel_size = 3, stride = 1,
89 padding = 1),
90             nn.BatchNorm2d(out_ch),
91             nn.ReLU())
92         self.conv2 = nn.Sequential(
93             nn.Conv2d(out_ch, out_ch, kernel_size = 3, stride =
94 1, padding = 1),
95             nn.BatchNorm2d(out_ch))
96         self.relu = nn.ReLU()
97
98         if downsample:
99             self.downsampler = nn.Conv2d(in_ch, out_ch, 1, stride=2)
100
101     def forward(self, x):
102         residual = x
103         out = self.conv1(x)
104
105         if self.in_ch == self.out_ch:
106             out = self.conv1(out)
107
108         if self.downsample:
109             out = self.downsampler(out)
110             residual = self.downsampler(residual)
111
112         if self.skip_connections:
113             if self.in_ch == self.out_ch:

```



```

111         out = out + residual
112     else:
113         # Assuming equivalent dimensions which this dataset fits
114         firstSection = out[:, :self.in_ch, :, :]
115         secondSection = out[:, self.in_ch: :, :]
116
117         out = torch.cat((firstSection + residual, secondSection +
118 residual), dim=1)
119
120     return out
121
122 # %% [markdown]
123 # # CNN
124
125 # %%
126 # CNN model # Inspired by DLStudio LOADnet2 && notes from class
127 class CNN(nn.Module):
128     def __init__(self, skip_connections=True, depth=8):
129         super(CNN, self).__init__()
130         self.skip_connections = skip_connections
131         self.depth = depth // 2
132
133         # Create base layers
134         self.conv = nn.Conv2d(in_channels=3, out_channels=64, kernel_size=3,
135 padding=1)
136         self.pool = nn.MaxPool2d(kernel_size=2, stride=2)
137
138         # Classification
139         self.bn1 = nn.BatchNorm2d(num_features=64)
140         self.bn2 = nn.BatchNorm2d(num_features=128)
141
142         # Create Depth layers 64
143         self.skip64_arr = nn.ModuleList()
144         for idx in range(self.depth):
145             self.skip64_arr.append(Block(in_ch=64, out_ch=64, skip_connections=
146 self.skip_connections))
147
148         self.skip64ds = Block(in_ch=64, out_ch=64, downsample=True,
149 skip_connections=self.skip_connections)
150         self.skip64to128 = Block(in_ch=64, out_ch=128, skip_connections=self.
151 skip_connections)
152
153         # Create depth layer 128
154         self.skip128_arr = nn.ModuleList()
155         for idx in range(self.depth):
156             self.skip128_arr.append(Block(in_ch=128, out_ch=128,
157 skip_connections=self.skip_connections))
158         self.skip128ds = Block(in_ch=128, out_ch=128, downsample=True,
159 skip_connections=self.skip_connections)
160
161         # Linear Layers #Limited from original file do to lack of memory space
162         self.fc1 = nn.Linear(in_features=32*32*128, out_features=3) # 3
163 categories
164
165         ## For Regression
166         self.conv_sequential = nn.Sequential(nn.Conv2d(in_channels=64,
167 out_channels=64, kernel_size=3, padding=1),

```

```

159         nn.BatchNorm2d(num_features=64),
160         nn.ReLU(inplace=True),
161         nn.Conv2d(in_channels=64, out_channels
=64, kernel_size=3, padding=1),
162         nn.ReLU(inplace=True))
163
164     self.fc_sequential = nn.Sequential(nn.Linear(in_features=128*128*64,
out_features=4)) # 4 dimensions bbox
165
166     def forward(self, x):
167         x = self.pool(nn.functional.relu(self.conv(x)))
168
169         # Classification
170         cls = x.clone()
171         for idx, skip64 in enumerate(self.skip64_arr[:self.depth//4]):
172             cls = skip64(cls)
173         cls = self.skip64ds(cls)
174         for idx, skip64 in enumerate(self.skip64_arr[self.depth//4:]):
175             cls = skip64(cls)
176         cls = self.bn1(cls)
177         cls = self.skip64to128(cls)
178         for idx, skip128 in enumerate(self.skip128_arr[:self.depth//4]):
179             cls = skip128(cls)
180         cls = self.bn2(cls)
181         cls = self.skip128ds(cls)
182         for idx, skip128 in enumerate(self.skip128_arr[self.depth//4:]):
183             cls = skip128(cls)
184         cls = cls.view(-1, 32 * 32 * 128)
185         cls = self.fc1(cls)
186
187         # Regression for BBox
188         bbox = self.conv_sequential(x)
189         bbox = bbox.view(x.size(0), -1)
190         bbox = self.fc_sequential(bbox)
191
192         return cls, bbox
193
194 # %% [markdown]
195 # # Training
196
197 # %%
198 ##### TRAINING #####
199 # Inspired by DLStudio run_code_for_training_with_CrossEntropy_and_MSE_Losses &&
notes from class
200 def trainMSERegression(model, trainDataLoader, runn_avg_size, epochs, lr=1e-4,
betas=(0.9, 0.99)):
201     model = model.to(device)
202     num_layers = len(list(model.parameters()))
203     print(f"Number of layers: {0}".format(num_layers))
204
205     classification_criterion = nn.CrossEntropyLoss()
206     regression_criterion = nn.MSELoss()
207
208     optimizer = torch.optim.Adam(model.parameters(), lr=lr, betas=betas)
209     labeling_loss_running_avg = []
210     regression_loss_running_avg = []
211

```

```

212     for epoch in range(1, epochs+1):
213         running_loss_labeling = 0.0
214         running_loss_regression = 0.0
215         for batch_idx, (imgTensor, labels, bbox) in enumerate(trainDataLoader):
216             imgTensor = imgTensor.to(device)
217             labels = labels.to(device)
218             bbox = bbox.to(device)
219
220             optimizer.zero_grad()
221             outputs = model(imgTensor)
222             output_label = outputs[0]
223             output_bbox = outputs[1]
224
225             loss_label = classification_criterion(output_label, labels)
226             running_loss_labeling += loss_label.item()
227             loss_label.backward(retain_graph=True)
228
229             loss_bbox = regression_criterion(output_bbox, bbox)
230             running_loss_regression += loss_bbox.item()
231             loss_bbox.backward()
232
233             optimizer.step()
234
235             if(batch_idx % runn_avg_size == (runn_avg_size - 1)):
236                 labeling_loss_running_avg.append(running_loss_labeling / float(
runn_avg_size))
237                 regression_loss_running_avg.append(running_loss_regression /
float(runn_avg_size))
238
239                 running_loss_labeling = 0.0
240                 running_loss_regression = 0.0
241
242         return labeling_loss_running_avg, regression_loss_running_avg
243
244 # %%
245 ##### TRAINING #####
246 # Inspired by DLStudio run_code_for_training_with_CrossEntropy_and_MSE_Losses &&
notes from class
247 def trainCompleteBoxIOU(model, trainDataLoader, runn_avg_size, epochs, lr=1e-4,
betas=(0.9, 0.99)):
248     model = model.to(device)
249     num_layers = len(list(model.parameters()))
250     print(f"Number of layers: {0}".format(num_layers))
251
252     classification_criterion = nn.CrossEntropyLoss()
253     regression_criterion = complete_box_iou_loss
254
255     optimizer = torch.optim.Adam(model.parameters(), lr=lr, betas=betas)
256     labeling_loss_running_avg = []
257     regression_loss_running_avg = []
258
259     for epoch in range(1, epochs+1):
260         running_loss_labeling = 0.0
261         running_loss_regression = 0.0
262         for batch_idx, (imgTensor, labels, bbox) in enumerate(trainDataLoader):
263             imgTensor = imgTensor.to(device)
264             labels = labels.to(device)

```

```

265         bbox = bbox.to(device)
266
267         optimizer.zero_grad()
268         outputs = model(imgTensor)
269         output_label = outputs[0]
270         output_bbox = outputs[1]
271
272         loss_label = classification_criterion(output_label, labels)
273         loss_label.backward(retain_graph=True)
274         running_loss_labeling += loss_label.item()
275
276         loss_bbox = regression_criterion(output_bbox, bbox, reduction = "
mean")
277         loss_bbox.backward()
278         running_loss_regression += loss_bbox.item()
279
280         optimizer.step()
281
282         if(batch_idx % runn_avg_size == (runn_avg_size - 1)):
283             labeling_loss_running_avg.append(running_loss_labeling / float(
runn_avg_size))
284             regression_loss_running_avg.append(running_loss_regression /
float(runn_avg_size))
285
286             running_loss_labeling = 0.0
287             running_loss_regression = 0.0
288
289         return labeling_loss_running_avg, regression_loss_running_avg
290
291 # %% [markdown]
292 # # Testing
293
294 # %%
295 def test(model, testDataLoader, class_list):
296     model = model.to(device)
297     confusion_matrix = np.zeros((len(class_list), len(class_list)))
298     image_data = []
299
300     with torch.no_grad():
301         for imgTensor, labels, bbox in testDataLoader:
302             # Set to Device
303             imgTensor = imgTensor.to(device)
304             labels = labels.to(device)
305             bbox = bbox.to(device)
306
307             # Set Outputs
308             outputs = model(imgTensor)
309             output_label = outputs[0]
310             output_bbox = outputs[1].tolist()
311
312             _, predicted = torch.max(output_label, dim=1)
313             for label, prediction in zip(labels, predicted):
314                 confusion_matrix[label][prediction] += 1
315
316             for img, original_label, predicted_label, original_bbox,
predicted_bbox in zip(imgTensor, labels, predicted, bbox, output_bbox):

```

```

317         image_data.append([img, original_label, predicted_label,
original_bbox, predicted_bbox])
318
319     accuracy = np.trace(confusion_matrix) / np.sum(confusion_matrix)
320     return confusion_matrix, accuracy, image_data

```

Listing 2: Skip-block CNN training and testing code

## 5 Task 3: Validation

For validation a loss plot vs iteration was made for the Cross-Entropy, MSE, and CIoU loss. A confusion matrix was further created to compare accuracy based on MSE and CIoU loss.

```

1  # %%
2  def plotRegressionLosses(regression, lossType):
3      figure = plt.figure(1)
4
5      plt.plot(range(len(regression)), regression, label="Loss")
6
7      plt.xlabel("Iterations")
8      plt.ylabel("Loss")
9      plt.legend(loc="lower right")
10
11     plt.savefig("lossvsiter_regress{0}.jpg".format(lossType))
12
13  # %%
14  def plotLabelLosses(labeling, lossType):
15      figure = plt.figure(2)
16
17      plt.plot(range(len(labeling)), labeling, label="Loss")
18
19      plt.xlabel("Iterations")
20      plt.ylabel("Loss")
21      plt.legend(loc="lower right")
22
23      plt.savefig("lossvsiter_class{0}.jpg".format(lossType))
24
25  # %%
26  def plotConfusionMatrix(conf, accuracy, class_list, lossType):
27      figure = plt.figure(3)
28      sns.heatmap(conf, xticklabels=class_list, yticklabels=class_list, annot=True)
29
30      plt.xlabel("True Label: Accuracy {0}".format(accuracy))
31      plt.ylabel("Predicted Label")
32      plt.savefig("confMatrix{0}.jpg".format(lossType))
33
34  # %%
35  def plotImages(image_data, lossType):
36      fig, ax = plt.subplots(3, 3)
37      row, col = 0, 0
38
39      for idx, arr in enumerate(image_data):
40          image = tvn.ToPILImage()(arr[0])
41
42          invert_categories = { # Set Integer Values for Cat
0 : "bus",

```

```

43         1 : "cat",
44         2 : "pizza"
45     }
46
47     original_cat = invert_categories[arr[1].item()]
48     predicted_cat = invert_categories[arr[2].item()]
49     original_bbox = arr[3]
50     predicted_bbox = arr[4]
51
52     # Original bbox
53     #print(original_bbox)
54     [x, y, w, h] = original_bbox
55     image = np.uint8(image)
56     image = cv2.rectangle(image, (int(x), int(y)), (int(x + w), int(y + h)),
57                             color=(36, 255, 12), thickness=2)
58     image = cv2.putText(image, original_cat, (int(x), int(y - 10)), fontFace
59 =cv2.FONT_HERSHEY_SIMPLEX, fontScale=0.8, color=(36, 255, 12), thickness=2)
60
61     [x, y, w, h] = predicted_bbox
62     # image = np.uint8(image)
63     image = cv2.rectangle(image, (int(x), int(y)), (int(x + w), int(y + h)),
64                             color=(255, 36, 12), thickness=2)
65     image = cv2.putText(image, predicted_cat, (int(x), int(y - 10)),
66                             fontFace=cv2.FONT_HERSHEY_SIMPLEX, fontScale=0.8, color=(255, 36, 12),
67                             thickness=2)
68
69     print(row, col)
70     ax[row, col].imshow(image)
71
72     # Increment through row
73     col += 1
74     if col == 3:
75         col = 0
76         row += 1
77
78     if idx > 7:
79         break
80
81     # fig, axs = plt.subplots(nrows=3, ncols=3, figsize=(10,10), dpi = 150)
82     # plt.show()
83     # axs = axs.flatten()
84     plt.savefig("predvsoriginal_bbox{0}.jpg".format(lossType))

```

**Listing 3:** *Plotting and Validation Code*

```

1 # Images saved directory
2 train_data_path = 'trainingData'
3 val_data_path = 'valData'
4
5 class_list = ['pizza', 'bus', 'cat']
6
7 trainDataFrame = pd.read_csv('train_labels.csv')
8 valDataFrame = pd.read_csv('val_labels.csv')
9
10 # %%
11 trainDataset = MyDataset(trainDataFrame, class_list)
12 valDataset = MyDataset(valDataFrame, class_list)

```

```

13
14 trainDataloader = torch.utils.data.DataLoader(trainDataset, batch_size=2,
    num_workers=2, shuffle=True)
15 valDataloader = torch.utils.data.DataLoader(valDataset, batch_size=2,
    num_workers=2, shuffle=True)
16
17 # %%
18 # Save networks
19 net = CNN()
20
21 # Train Networks
22 epochs=7
23 labelMSELoss, regressMSELosses = trainMSERegression(net, trainDataloader, 100,
    epochs=epochs, lr=1e-4, betas=(0.9, 0.99))
24 lossType = 'MSELoss'
25 plotRegressionLosses(regressMSELosses, lossType)
26 plotLabelLosses(labelMSELoss, lossType)
27
28 conf, acc, image_data = test(net, valDataloader, class_list)
29 plotConfusionMatrix(conf, acc, class_list, lossType)
30 plotImages(image_data, lossType)
31
32 labelCBIoULoss, regressCBIoULosses = trainCompleteBoxIOU(net, trainDataloader,
    100, epochs=epochs, lr=1e-4, betas=(0.9, 0.99))
33 lossType = "CIoULoss"
34 plotRegressionLosses(regressCBIoULosses, lossType)
35 plotLabelLosses(labelCBIoULoss, lossType)
36
37 conf, acc, image_data = test(net, valDataloader, class_list)
38 plotConfusionMatrix(conf, acc, class_list, lossType)
39 plotImages(image_data, lossType)

```

**Listing 4:** *Main for running all prior code*


The following images are the results and validation for the model created in the prior code. It will be shown that all models met criteria of high enough accuracy and that the number of layer the system surpassed is greater than 50 as seen in Figure 2.




Number of layers: 108

**Figure 2:** *Layer in network*

The model was made up of 108 layers that, when trained with MSE, granted a accuracy of 0.81 and, when trained with CIoU, granted a accuracy of 0.8. This indicates that MSE Regression Loss outperforms CIoU but this could be due to the effects of one being run a cluster with gpu access and another run on google colab. Based on the actual figure 6 and 7 one can see that MSE performs better than the other in terms of fitting boxes which seems to be due to a confusion another object in the image. The network could be improved by simply adding more layers, running for more epochs, or shifting the learning as with any model. A unique solution would be changing the skip-block made to resemble more advance methods that do not make the same assumptions that our does.




(a) *Cross-Entropy Loss Per Iteration*



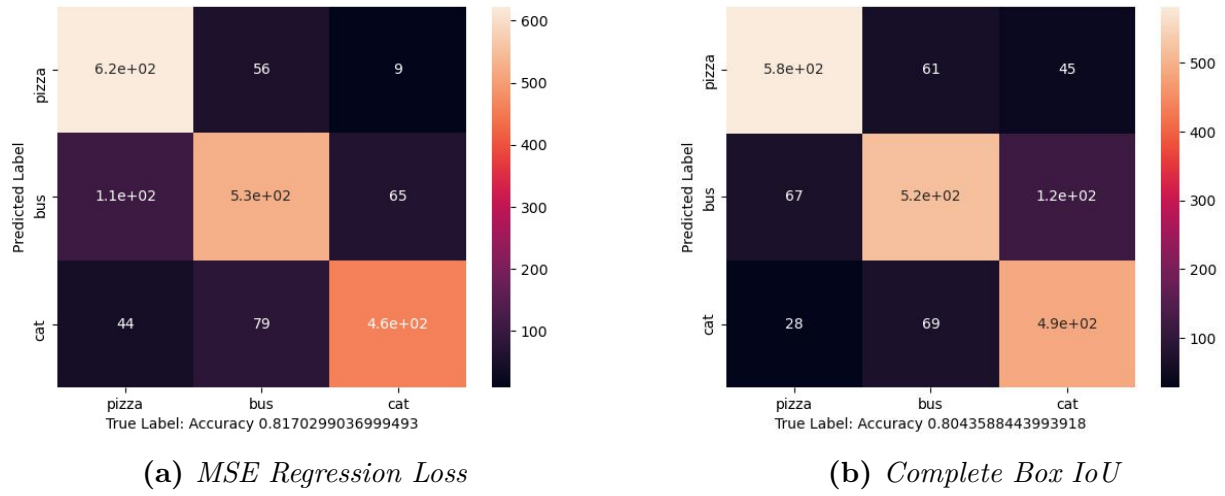
(b) *MSE Regression Loss*

**Figure 3:** *Training Loss for Cross-Entropy and MSE loss*

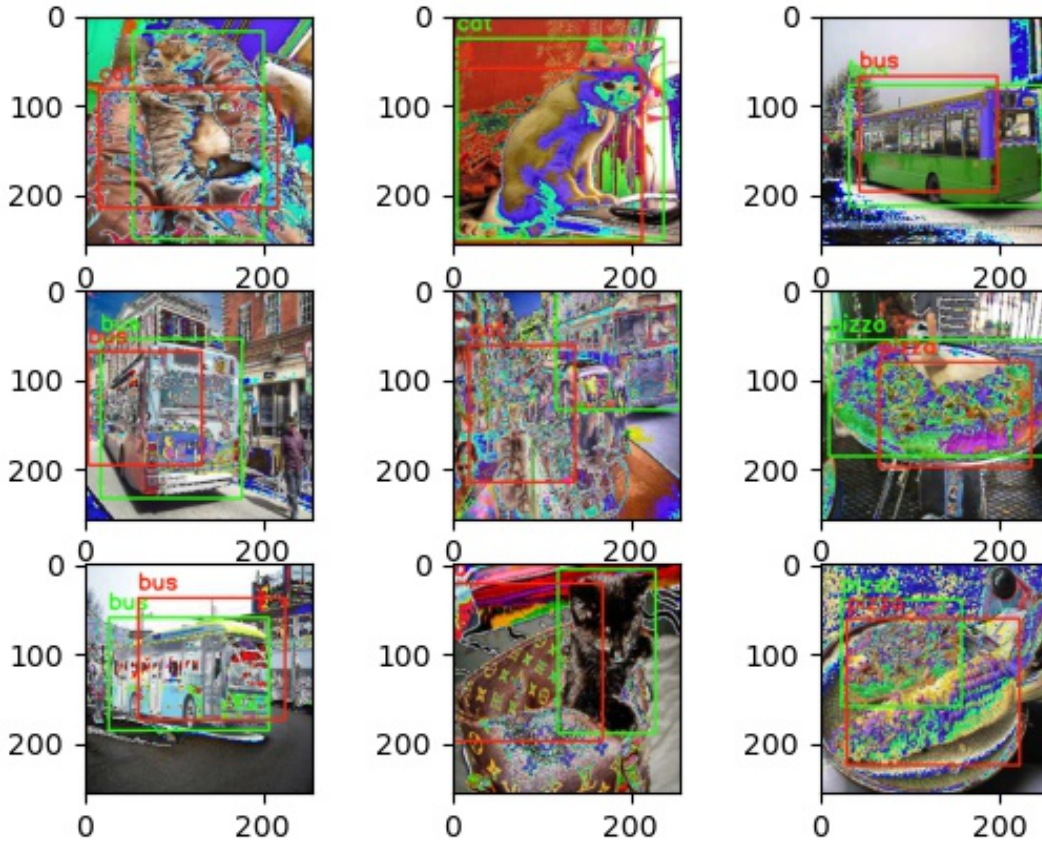


**Figure 4:** *Complete Box IoU Loss*

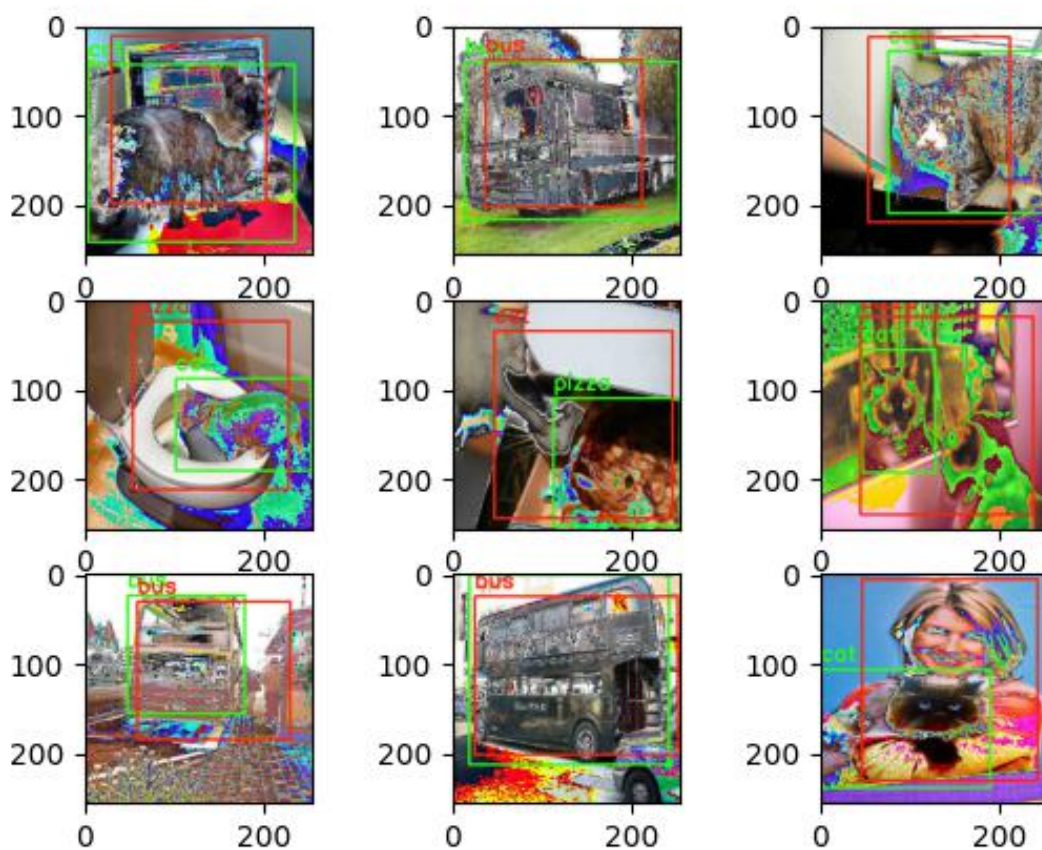




**Figure 5:** *Confusion Matrixes*



**Figure 6:** *MSE Regression Loss bounding box results*



**Figure 7:** *CIoU Regression Loss bounding box results*

## 6 Lessons learned

From this assignment, I learned about boundary boxes, skip-block implementation, and CIoU.