

# HW 6: YOLO

David E. Farache, Email ID: dfarache@purdue.edu

March 28, 2023

## 1 Introduction

There are different methods for multi-instance processing of images and their boundary boxes, including R-CNN, SSD, and YOLO. This work focuses on "You Only Look Once" or the YOLO method, into a convolutional neural network (CNN). This work should implement skip-blocks created from the previous homework and using Binary Cross-Entropy (BCE) loss, Mean Squared Error (MSE) loss, and Cross-Entropy (CE) loss for training.

## 2 Methodology

COCO was utilized to create an initial training set via the `pycocotools.coco` function, which would reduce to 6890 for training and 3494 for validation. The data-loader was based on HW4 with the addition of encoding based on the category and transform the image from a PIL image to a normalized tensor.

YOLO method functions by dividing the image into a grid system and the prediction of the boundary box and the class-label is designated to the grid-cell whereby the center of the image is. This enables for single network to be trained for multiple labels as the separation is made with the grid-cells and anchor boxes, instead of multiple networks as R-CNN would require. For loss calculation, the existence of the image within the cell is done using BCE, classification is measured using CE, and refinement and decision of the bounding box is based upon the MSE regression loss, using the equation seen below.

$$BCE = -(y \log(p) + (1 - y) \log(1 - p)) \quad (1)$$

$$MSE = \sum_{i=1}^D (x_i - y_i)^2 \quad (2)$$

$$Cross - Entropy = - \sum_{c=1}^M y_{o,c} \log(p_{o,c}) \quad (3)$$

## 3 Task 1: COCO Data-set Preparation

For task 1, the COCO data-set was used from train2014 and val2014 in order to reduce and create a training and validation for categories: pizza, cat, and bus. Images were selected based on whether there was at least one dominant label with area greater than 64 x 64. Each image was turned in RGB and resized from their original height to 256 x 256. All bounding boxes of the images were

re-scaled and there accompanying category saved to a list. A data-frame of filename, category list, and boundry box list, were saved for efficient processing.

```
1 # %% [markdown]
2 # # Library
3
4 # %%
5 %matplotlib inline
6 from pycocotools.coco import COCO
7
8 import numpy as np
9 import skimage.io as io
10 import skimage
11 import cv2
12 import pandas as pd
13
14 import matplotlib.pyplot as plt
15 import pylab
16 import random
17 from PIL import Image
18 pylab.rcParams['figure.figsize'] = (8.0, 10.0)
19
20 # %% [markdown]
21 # # Display Images
22
23 # %%
24 inverse_categories = {6: "bus", 17: "cat", 59: "pizza"}
25
26 # %%
27 # Based from skeleton code given
28 def display_random_image_with_bbox(image_path, bboxes, cat_set):
29     file = image_path
30     image = Image.open(file)
31
32     print(bboxes)
33     print(cat_set)
34     for bbox, cat in zip(bboxes, cat_set):
35         [x, y, w, h] = bbox
36         fig, ax = plt.subplots(1,1)
37         image = np.uint8(image)
38         image = cv2.rectangle(image, (int(x), int(y)), (int(x + w), int(y + h)),
39                               color=(36, 255, 12), thickness=2)
40         image = cv2.putText(image, inverse_categories[cat], (int(x), int(y - 10)
41 ), fontFace=cv2.FONT_HERSHEY_SIMPLEX, fontScale=0.8, color=(36, 255, 12),
42 thickness=2)
43
44         ax.imshow(image)
45         ax.set_axis_off()
46     plt.axis("tight")
47     plt.show()
48
49     fig, axs = plt.subplots(nrows=3, ncols=3, figsize=(10,10), dpi = 150)
50     axs = axs.flatten()
51     axs_count = 0
52
53     # %%
54     def display_images(saveForPlotting):
```

```

52 fig, ax = plt.subplots(3, 3)
53 row, col = 0, 0
54
55 #print(len(saveForPlotting))
56 for arr in saveForPlotting:
57     file = arr[0]
58     image = Image.open(file)
59     print(arr)
60     for cat, bbox in zip(arr[2], arr[1]):
61         [x, y, w, h] = bbox
62         image = np.uint8(image)
63         image = cv2.rectangle(image, (int(x), int(y)), (int(x + w), int(y +
h)), color=(36, 255, 12), thickness=2)
64         image = cv2.putText(image, inverse_categories[cat], (int(x), int(y -
10)), fontFace=cv2.FONT_HERSHEY_SIMPLEX, fontScale=0.8, color=(36, 255, 12),
thickness=2)
65         ax[row, col].imshow(image)
66
67     # Increment through row
68     col += 1
69
70     if col == 3:
71         col = 0
72         row += 1
73
74 fig, axes = plt.subplots(nrows=3, ncols=3, figsize=(10,10), dpi = 150)
75 plt.show()
76 axes = axes.flatten()
77
78 # %% [markdown]
79 # # Adjust and Change Images
80
81 # %%
82 def resizeBBOX(bbox, startWidth, startHeight, resize):
83     x_scale = resize / startWidth
84     y_scale = resize / startHeight
85     bboxResize = np.zeros(4)
86
87     bboxResize[0] = int(x_scale * bbox[0])
88     bboxResize[1] = int(y_scale * bbox[1])
89     bboxResize[2] = int(x_scale * bbox[2])
90     bboxResize[3] = int(y_scale * bbox[3])
91
92     return bboxResize.tolist()
93
94 # %%
95 def resizeAndRepairImage(start_image_path, new_image_path, filename, resize):
96     img = Image.open(start_image_path + '/' + filename)
97     width, height = img.size
98
99     image = img.convert(mode="RGB")
100
101     img = image.resize((resize, resize), Image.BOX)
102     img.save(new_image_path + '/' + filename)
103
104 # %%
105 def saveAsDataFrame(ids, cats, filepaths, x1, y1, width, height, type):

```

```

106 # Containers for data to go in training label csv
107 columns = ["id", "category", "filepath", "x1", "y1", "width", "height"]
108 dataFrame = pd.DataFrame(columns=columns)
109 dataFrame = dataFrame.astype('object')
110
111 dataFrame["id"] = np.array(ids)
112 dataFrame["category"] = cats
113 dataFrame["filepath"] = filepaths
114 dataFrame["x1"] = np.array(x1)
115 dataFrame["y1"] = y1
116 dataFrame["width"] = width
117 dataFrame["height"] = height
118
119 dataFrame.to_csv("{0}_labels.csv".format(type))
120
121 return dataFrame
122
123
124 # %% [markdown]
125 # # Choose Images Within Parameters
126
127 # %%
128 # Get Random Images from set
129 def ImageSelection(start_image_path, new_image_path, cocoObj, class_list, type,
spotSave):
130     # Save File Location List
131     saveImportant = []
132     saveForPlotting = []
133
134     # Save image info
135     ids = []
136     cats = []
137     filepaths = []
138     x1 = []
139     y1 = []
140     widthFrame = []
141     heightFrame = []
142     for cat in class_list:
143         # get all images containing given categories
144         catIds = cocoObj.getCatIds(catNms=[cat]) # Get ids from annotations
145         imgIds = cocoObj.getImgIds(catIds=catIds) # Load images ids of chosen
146         annotations ids
147         img = cocoObj.loadImgs(ids=imgIds) # Get images
148         numPlots = 0 # Make X number plots
149
150         #Loop per image
151         for idx, images in enumerate(img):
152             annIds = cocoObj.getAnnIds(imgIds=images['id'], catIds=catIds,
iscrowd=False) # Get dictionary value
153             anns = cocoObj.loadAnns(annIds) # Get annotations
154
155             # Saving List
156             id_set = []
157             cat_set = []
158             x1_set = []
159             y1_set = []
160             widthFrame_set = []

```

```

160         heightFrame_set = []
161         bbox_set = []
162
163         forGroundImge = 0
164         for jdx, ann in enumerate(anns):
165
166             width, height = images['width'], images['height']
167
168             if ann['area'] > 64 * 64: #check if means parameter of dominate
obj
169                 # Resize Images
170                 bboxResize = resizeBBBOX(ann['bbox'], int(width), int(height)
, resize=256) #Adjust Box
171
172                 # Append info
173                 id_set.append(ann['id'])
174                 cat_set.append(ann['category_id'])
175                 x1_set.append(bboxResize[0])
176                 y1_set.append(bboxResize[1])
177                 widthFrame_set.append(bboxResize[2])
178                 heightFrame_set.append(bboxResize[3])
179                 bbox_set.append(bboxResize)
180
181                 forGroundImge = 1
182
183             if forGroundImge == 1:
184                 resizeAndRepairImage(start_image_path, new_image_path, images['
file_name'], resize=256) # Save Adjust Image
185
186                 cats.append(cat_set)
187                 filepaths.append("{0}/{1}".format(spotSave, images['file_name']))
188                 ids.append(id_set)
189                 x1.append(x1_set)
190                 y1.append(y1_set)
191                 widthFrame.append(widthFrame_set)
192                 heightFrame.append(heightFrame_set)
193
194                 if numPlots < 3 and len(cat_set) > 1:
195                     #display_random_image_with_bbox("{0}/{1}".format(spotSave,
images['file_name']), bbox_set, cat_set)
196                     saveForPlotting.append(["{0}/{1}".format(spotSave, images['
file_name']), bbox_set, cat_set])
197                     numPlots += 1
198
199                 display_images(saveForPlotting)
200                 dataframe = saveAsDataFrame(ids, cats, filepaths, x1, y1, widthFrame,
heightFrame, type)
201                 return dataframe
202
203 # %%
204 # Input
205 train_json = '/Users/davidfarache/Documents/ECE60146/HW6/annotations/
instances_train2014.json'
206 val_json = '/Users/davidfarache/Documents/ECE60146/HW6/annotations/
instances_val2014.json'
207
208 train_path = '/Users/davidfarache/Documents/ECE60146/HW6/train2014'

```

```

209 train_data_path = '/Users/davidfarache/Documents/ECE60146/HW6/trainingData'
210
211 val_path = '/Users/davidfarache/Documents/ECE60146/HW6/val2014'
212 val_data_path = '/Users/davidfarache/Documents/ECE60146/HW6/valData'
213
214 trainSaveSpot = 'trainingData'
215 valSaveSpot = 'valData'
216
217 class_list = ['pizza', 'bus', 'cat']
218
219 # %%
220 cocoTrain = COCO(train_json)
221 cocoVal = COCO(val_json)
222
223 trainDataFrame = ImageSelection(train_path, train_data_path, cocoTrain,
                                class_list, 'train', trainSaveSpot)
224 valDataFrame = ImageSelection(val_path, val_data_path, cocoVal, class_list, 'val',
                                valSaveSpot)
225
226 print(len(trainDataFrame))
227 print(len(valDataFrame))
228
229 # %%
230 print(trainDataFrame.columns)
231 print(valDataFrame.columns)
232
233 # %%
234 #print(trainDataFrame['category'])
235 saving_set = []
236 for i, val in enumerate(trainDataFrame['category']):
237     if val[0] not in saving_set:
238         print(trainDataFrame['filepath'].iloc[i])
239         saving_set.append(val[0])
240         print(val[0])

```

**Listing 1:** *Creating COCO data-set and plotting*

A sample of the images select has been placed below, which can be seen to have shrunk and had their boundary box rescaled within them:



Figure 1: *Selection of Training Set*

## 4 Task 2: Create YOLO Vector

The following code has the data set that stores the boundary boxes, normalized images, and image labels. The code below is what is required to create the data-loader which includes the YOLO tensor. The YOLO tensor code takes in the boundary boxes, finds the center of the x,y cord for the boundary box image, gets the height and width of the box, and the difference of distance between a YOLO cell center and its own to establish which is responsible. The aspect ratio is taken to set which anchor box is saved and all information is saved as a tensor of  $[1, \delta_x, \delta_y, h, w, n\text{-categories}]$  for one hot encoding (in this case 3). In this case the image grid was selected to be 6x6 meaning each must have 42x42 pixels. The aspect ratios calculated are those given in the lecture: 1/5, 1/3, 1/1, 3/1, and 5/1.

```

1 # Import Libraries
2 import numpy as np
3 import torch
4 import torchvision.transforms as tvt
5 import torch.utils.data
6 import torch.nn as nn

```

```

7 import torch.nn.functional as F
8 import matplotlib.pyplot as plt
9 from PIL import Image
10 import os
11 import seaborn as sns
12 from torchvision.ops import box_convert
13 import pandas as pd
14 import cv2
15 import json
16 from torchinfo import summary
17 from copy import deepcopy
18
19 # GLOBAL VARIABLES
20 device = 'cuda'
21 device = torch.device(device)
22
23 # For DataLoader
24 root = '/scratch/gilbreth/dfarache/ece60146/David/HW6/'
25 class_list = ["bus", "cat", "pizza"]
26 inverse_categories = {6: "bus", 17: "cat", 59: "pizza"}
27
28 # Global YOLO Values (had to separate due to no longer sharing training and yolo
    vecotr in same function)
29 num_anchor_boxes = 5 # (height/width) 1/5 1/3 1/1 3/1 5/1 aspect ratios
30 max_num_objects = len(class_list)
31
32 yolo_vector_size = 8 # []
33 yolo_interval = 42 # Each cell is 42x42 pixels
34
35 num_cells_image_width = 256 // yolo_interval
36 num_cells_image_height = 256 // yolo_interval
37
38 num_yolo_cells = num_cells_image_width * num_cells_image_height # Create a grid
    of cells overlaying the image
39
40 # YOLO Tensor
41
42 # Based on DLStudio run_code_for_training_multi_instance_detection function and
    homework example
43 def createYoloTensor(bboxes, labels, num_images_in_batch=1):
44     yolo_tensor = torch.zeros(num_yolo_cells, num_anchor_boxes, yolo_vector_size
    )
45
46     # Create empty torch tensors
47     height_center_bb = torch.zeros(num_images_in_batch, 1).float()
48     width_center_bb = torch.zeros(num_images_in_batch, 1).float()
49     object_bb_height = torch.zeros(num_images_in_batch, 1).float()
50     object_bb_width = torch.zeros(num_images_in_batch, 1).float()
51
52     numericDict = {6: 0, 17: 1, 59: 2} # Swap id to index for one-hot encoding
53
54     # i is index of object in the foreground
55     for i in range(max_num_objects):
56         # remove .float
57         y_cord_center = (bboxes[i, 1] + bboxes[i, 3] / 2.0).int() # Get y-
        coordinate object center from y1

```



```

58     x_cord_center = (bboxes[i, 0] + bboxes[i, 2] / 2.0).int() # Get x-
coordinate object center from x1
59
60     object_bb_height = (bboxes[i, 3] - bboxes[i, 1]).float() # Height bounding
box
61     object_bb_width = (bboxes[i, 2] - bboxes[i, 0]).float() # Width bounding
box
62
63     if (object_bb_height < 4.0) or (object_bb_width < 4.0): continue
64
65     # Get the cell row and column index that corresponds to the center of
the bounding box
66     cell_row_i = torch.clamp((y_cord_center / yolo_interval).int(), max=
num_cells_image_height - 1)
67     cell_col_i = torch.clamp((x_cord_center / yolo_interval).int(), max=
num_cells_image_width - 1)
68
69     # Get the height of the bounding box divided by the actual height of the
cell
70     bheight = y_cord_center / yolo_interval
71     bwidth = x_cord_center / yolo_interval
72
73     # Swap from x,y system to i,j coordinate
74     cell_center_i = cell_row_i * yolo_interval + float(yolo_interval) / 2.0
75     cell_center_j = cell_col_i * yolo_interval + float(yolo_interval) / 2.0
76
77     # Compute del_x and del_y
78     del_x = (x_cord_center.float() - cell_center_j.float()) / yolo_interval
79     del_y = (y_cord_center.float() - cell_center_i.float()) / yolo_interval
80
81     # Get the class label
82     class_label_of_object = int(labels[i].item())
83     if class_label_of_object == 31: continue # Disregard labels with class
label 31
84
85     # Get Aspect Ratio
86     aspect_ratio = object_bb_height / object_bb_width
87     anchor_box_idx = 0
88     if aspect_ratio <= 0.2:
89         anchor_box_idx = 0
90     if 0.2 < aspect_ratio <= 0.5:
91         anchor_box_idx = 1
92     if 0.5 < aspect_ratio <= 1.5:
93         anchor_box_idx = 2
94     if 1.5 < aspect_ratio <= 4.0:
95         anchor_box_idx = 3
96     if aspect_ratio > 4.0:
97         anchor_box_idx = 4
98
99     # Create the yolo vector
100    # Vector [exsit, x1, y1, height, width, n-encoding]
101    yolo_vector = torch.FloatTensor([1, del_x.item(), del_y.item(), bheight.
item(), bwidth.item(), 0, 0, 0])
102
103    yolo_vector[5 + numericDict[class_label_of_object]] = 1 # One-hot
encoding
104
105    # Assign to yolo tensor

```

```

101         yolo_cell_index = cell_row_i.item() * num_cells_image_width + cell_col_i
102         .item()
103         yolo_tensor[yolo_cell_index, anchor_box_idx] = yolo_vector # place into
104         proper index
105     # Create an augmented yolo tensor
106     yolo_tensor_aug = torch.zeros(num_yolo_cells, num_anchor_boxes,
107     yolo_vector_size + 1).float()
108     yolo_tensor_aug[:, :, :-1] = yolo_tensor
109     # If not exist throw
110     for icx in range(num_yolo_cells):
111         for iax in range(num_anchor_boxes):
112             if(yolo_tensor_aug[icx, iax, 0] == 0):
113                 yolo_tensor_aug[icx, iax, -1] = 1
114
115     return yolo_tensor_aug
116
117 # Create DataLoader Code
118
119 def ImageProcessing(images, root):
120     # Get dir
121     image_dir = images["filepath"]
122
123     # Get inputs
124     image = Image.open(root + image_dir)
125
126     # Normalize image
127     toTensor = tvn.ToTensor()(image)
128     toNormalize = tvn.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))(toTensor)
129     return toNormalize
130
131 # Generate Datasets
132
133 class MyDataset(torch.utils.data.Dataset):
134     def __init__(self, imagesDataFrame, root, transform=None):
135         super().__init__()
136         self.imagesDataFrame = imagesDataFrame
137         self.root = root # Directory for folder images
138
139     def bboxAndLabels(self, imageInfo):
140         # Split bbox into parts
141         x1 = imageInfo["x1"]
142         y1 = imageInfo["y1"]
143         width = imageInfo["width"]
144         height = imageInfo["height"]
145
146         # Get Category
147         category = imageInfo["category"]
148
149         labels = torch.zeros(max_num_objects, dtype=torch.uint8) + 31
150         bboxes = torch.zeros(max_num_objects, 4, dtype=torch.uint8)
151
152         # Extract info from dataframe
153         x1_cat = x1[1:-1].split(',')

```

```

155     y1_cat = y1[1:-1].split(',')
156     width_cat = width[1:-1].split(',')
157     height_cat = height[1:-1].split(',')
158     labelList = category[1:-1].split(',')
159
160     #print(x1_cat, y1_cat, width_cat, height_cat, labelList)
161     for j, cat in enumerate(labelList):
162         create_box = [float(x1_cat[j]), float(y1_cat[j]), float(width_cat[j]
163         ), float(height_cat[j])]
164
165         if (j < max_num_objects):
166             bboxes[j] = torch.tensor(create_box, dtype=torch.float)
167             labels[j] = int(cat)
168
169     return labels, bboxes
170
171 def __len__(self):
172     return len(self.imagesDataFrame)
173
174 def __getitem__(self, i):
175     imageInfo = self.imagesDataFrame
176     imageInfo = imageInfo.iloc[i]
177     normalImage = ImageProcessing(imageInfo, self.root)
178
179     # Make YOLO Tensor for processing
180     labels, bboxes = self.bboxAndLabels(imageInfo)
181     yoloTensor = createYoloTensor(bboxes, labels)
182
183     return normalImage, bboxes, labels, yoloTensor

```

Listing 2: Setup YOLO Tensor and DataLoader

## 5 Task 3: Create Network and Train

The code below is based on that from hw 5, which grants 102 layers within the network.

```

1 # Based on DLStudio SkipBlock Code
2 # Based off following https://blog.paperspace.com/writing-resnet-from-scratch-in-
  -pytorch/
3 class Block(nn.Module):
4     def __init__(self, in_ch, out_ch, downsample=False, skip_connections=True):
5         super(Block, self).__init__()
6         self.downsample = downsample
7         self.in_ch = in_ch
8         self.out_ch = out_ch
9
10        self.skip_connections = skip_connections
11
12        self.conv1 = nn.Sequential(
13            nn.Conv2d(in_ch, out_ch, kernel_size = 3, stride = 1,
padding = 1),
14            nn.BatchNorm2d(out_ch),
15            nn.ReLU())
16        self.conv2 = nn.Sequential(
17            nn.Conv2d(out_ch, out_ch, kernel_size = 3, stride = 1,
padding = 1),

```

```

18         nn.BatchNorm2d(out_ch))
19     self.relu = nn.ReLU()
20
21     if downsample:
22         self.downsampler = nn.Conv2d(in_ch, out_ch, 1, stride=2)
23
24     def forward(self, x):
25         residual = x
26         out = self.conv1(x)
27
28         if self.in_ch == self.out_ch:
29             out = self.conv1(out)
30
31         if self.downsample:
32             out = self.downsampler(out)
33             residual = self.downsampler(residual)
34
35         if self.skip_connections:
36             if self.in_ch == self.out_ch:
37                 out = out + residual
38             else:
39                 # Assuming equivalent dimensions which this dataset fits
40                 firstSection = out[:, :self.in_ch, :, :]
41                 secondSection = out[:, self.in_ch :, :, :]
42
43                 out = torch.cat((firstSection + residual, secondSection +
44 residual), dim=1)
45
46         return out
47
48 # CNN model
49 # Based on DLStudio LOADnet2 && notes from class
50 class NetForYolo(nn.Module):
51     # Inspired by Professor Kak's NetForYolo class
52     def __init__(self, skip_connections=True, depth=8):
53         super(NetForYolo, self).__init__()
54         self.skip_connections = skip_connections
55         self.depth = depth // 2
56
57         self.conv1 = nn.Conv2d(in_channels=3, out_channels=64, kernel_size=3,
58 padding=1)
59         self.conv2 = nn.Conv2d(in_channels=64, out_channels=64, kernel_size=3,
60 padding=1)
61
62         self.pool = nn.MaxPool2d(kernel_size=2, stride=2)
63
64         self.bn1 = nn.BatchNorm2d(num_features=64)
65         self.bn2 = nn.BatchNorm2d(num_features=128)
66
67         self.skip64_arr = nn.ModuleList()
68         for idx in range(self.depth):
69             self.skip64_arr.append(Block(in_ch=64, out_ch=64, skip_connections=
70 self.skip_connections))
71
72         self.skip64ds = Block(in_ch=64, out_ch=64, downsample=True,
73 skip_connections=self.skip_connections)

```

```

69         self.skip64to128 = Block(in_ch=64, out_ch=128, skip_connections=self.
skip_connections)
70
71         self.skip128_arr = nn.ModuleList()
72         for idx in range(self.depth):
73             self.skip128_arr.append(Block(in_ch=128, out_ch=128,
skip_connections=self.skip_connections))
74
75         self.skip128ds = Block(in_ch=128, out_ch=128, downsample=True,
skip_connections=self.skip_connections)
76         self.skip128to256 = Block(in_ch=128, out_ch=256, skip_connections=self.
skip_connections)
77         self.skip256_arr = nn.ModuleList()
78         for idx in range(self.depth):
79             self.skip256_arr.append(Block(in_ch=256, out_ch=256,
skip_connections=self.skip_connections))
80
81         self.skip256ds = Block(in_ch=256, out_ch=256, downsample=True,
skip_connections=self.skip_connections)
82
83         self.fc_seqn = nn.Sequential(
84             nn.Linear(in_features=128*16*16, out_features=6*6*5*9),
85             nn.ReLU(inplace=True),
86             nn.Linear(in_features=4096, out_features=2048),
87             nn.ReLU(inplace=True),
88             nn.Linear(in_features=6*6*5*9, out_features=6*6*5*9) # 6
x6 grid
89         )
90     def forward(self, x):
91         x = self.pool(F.relu(self.conv1(x)))
92         x = self.pool(F.relu(self.conv2(x)))
93         for i, skip64 in enumerate(self.skip64_arr[:self.depth//4]):
94             x = skip64(x)
95         x = self.skip64ds(x)
96         for i, skip64 in enumerate(self.skip64_arr[:self.depth//4]):
97             x = skip64(x)
98         x = self.bn1(x)
99         x = self.skip64to128(x)
100        for i, skip128 in enumerate(self.skip128_arr[:self.depth//4]):
101            x = skip128(x)
102        x = self.bn2(x)
103        x = self.skip128ds(x)
104        x = x.view(-1, 128*16*16)
105        x = self.fc_seqn(x)
106        return x

```

**Listing 3:** *Make Network*

For training, the criterion of BCE, MSE, and CE was utilized. The inputs for the model were batch\_size 64, learning rate 1e-5, and trained for 8 epochs. Each YOLO vector is iterated through with values extracted for training and a running average of the losses taken every 25 values in the batch.

```

1 # Based on DLStudio run_code_for_training_multi_instance_detection function and
notes in class
2 def train(net, trainLoader, epochs, lr, betas):
3     print("Training")
4     optimizer = torch.optim.Adam(net.parameters(), lr=lr, betas=betas)

```

```

5
6 # Criteria
7 criterion1 = nn.BCELoss() # If object in image present
8 criterion2 = nn.MSELoss() # Regression for bounding box
9 criterion3 = nn.CrossEntropyLoss() # One hot encoding of label
10
11 net = net.to(device)
12 BCELossIter = []
13 MSELossIter = []
14 CELossIter = []
15
16 for epoch in range(1, epochs + 1):
17     # Initialize Loss Values
18     avgBCE_loss, avgMSE_loss, avgCE_loss = 0.0, 0.0, 0.0
19
20     for i, (images, _, _, yolo_tensors) in enumerate(trainLoader):
21         images = images.to(device)
22         yolo_tensors = yolo_tensors.to(device)
23
24         optimizer.zero_grad()
25
26         output = net(images)
27         output = output.view(batch_size, num_yolo_cells, num_anchor_boxes,
28                               yolo_vector_size+1)
29
30         totalBCE_loss = torch.tensor(0.0, requires_grad=True).float().to(
31             device)
32         totalMSE_loss = torch.tensor(0.0, requires_grad=True).float().to(
33             device)
34         totalCE_loss = torch.tensor(0.0, requires_grad=True).float().to(
35             device)
36
37         # Set yolo_tensor
38         yolo_input_tensor = torch.nonzero(output[:, :, :, 0])
39
40         # Seperate Values
41         batch_axis = yolo_input_tensor[:,0]
42         yolo_cells = yolo_input_tensor[:,1]
43         anchor_box = yolo_input_tensor[:,2]
44
45         # BCE Loss
46         bce_loss = criterion1(nn.Sigmoid()(output[:, :, :, 0]), yolo_tensors
47            [:, :, :, 0])
48         totalBCE_loss += bce_loss
49
50         # MSE Loss
51         predicted_regression_vector = output[batch_axis, yolo_cells,
52             anchor_box, 1:5]
53         target_regression_vector = yolo_tensors[batch_axis, yolo_cells,
54             anchor_box, 1:5]
55         mse_loss = criterion2(predicted_regression_vector,
56             target_regression_vector)
57         totalMSE_loss += mse_loss
58
59         # CELoss
60         class_probs_vector = output[batch_axis, yolo_cells, anchor_box,
61             5:-1]

```

```

53         target_class_vector = torch.argmax(yolo_tensors[batch_axis,
yolo_cells, anchor_box, 5:-1], dim=1)
54         ce_loss = criterion3(class_probs_vector, target_class_vector)
55         totalCE_loss += ce_loss
56
57         totalBCE_loss.backward(retain_graph=True)
58         totalMSE_loss.backward(retain_graph=True)
59         totalCE_loss.backward(retain_graph=True)
60
61         optimizer.step()
62
63         avgBCE_loss += totalBCE_loss.item()
64         avgMSE_loss += totalMSE_loss.item()
65         avgCE_loss += totalCE_loss.item()
66
67         numRunAvg = 25.0
68         if (i+1) % numRunAvg == 0:
69             BCELossIter.append(avgBCE_loss / numRunAvg)
70             MSELossIter.append(avgMSE_loss / numRunAvg)
71             CELossIter.append(avgCE_loss / numRunAvg)
72
73             print("{0} / {1}: {2}, {3}, {4}".format(epoch, epochs,
avgBCE_loss, avgMSE_loss, avgCE_loss))
74
75             avgBCE_loss, avgMSE_loss, avgCE_loss = 0.0, 0.0, 0.0
76
77     return BCELossIter, MSELossIter, CELossIter

```

Listing 4: *Training Code*

## 6 Task 4: Testing and Plotting

Plotting for loss over iteration and the resulting images predictions are shown below.

```

1 # Skeleton Code from HW5
2 def display_images(batch_images, batch_labels, batch_bboxes, predicted_labels,
predicted_bboxes):
3
4     fig, axes = plt.subplots(1)
5     batch_images = batch_images / 2 + 0.5
6     image = np.asarray(tvt.ToPILImage()(batch_images))
7     image_gt = np.asarray(tvt.ToPILImage()(batch_images))
8
9     # Labels
10    get_cat = [inverse_categories[int(i)] for i in batch_labels if int(i) != 31]
11    # 31
12
13    # Bounding Box
14    bboxes_in_image = [i for i in batch_bboxes if i.tolist() != [0,0,0,0]]
15
16    for j in range(len(bboxes_in_image)):
17        # Base
18        [x, y, w, h] = batch_bboxes[j]
19        image_gt = cv2.rectangle(image_gt, (int(x), int(y)), (int(w), int(h)),
color=(0,255,0), thickness=2)

```

```

19         image_gt = cv2.putText(image_gt, get_cat[j], ((int(x)), (int(y) - 10)),
fontFace=cv2.FONT_HERSHEY_SIMPLEX, fontScale=0.8, color=(0,255,0), thickness
=2)
20
21 #         axes[0].imshow(image_gt)
22
23     for j in range(len(predicted_bboxes)):
24         # Predicted
25         [x, y, w, h] = predicted_bboxes[j]
26         image_pred = cv2.rectangle(image, (int(x), int(y)), (int(w), int(h)),
color=(0,255,0), thickness=2)
27         image_pred = cv2.putText(image_pred, predicted_labels[j], ((int(x)), (
int(y) - 10)), fontFace=cv2.FONT_HERSHEY_SIMPLEX, fontScale=0.8, color
=(0,255,0), thickness=2)
28
29     axes.imshow(image_pred)

```

Listing 5: *Plotting*

```

1 # Testing Code
2
3 def AnalyzePred(pred_regress_vec, icx):
4     del_x, del_y = pred_regress_vec[0], pred_regress_vec[1]
5     h, w = pred_regress_vec[2], pred_regress_vec[3]
6
7     h *= yolo_interval
8     w *= yolo_interval
9
10    cell_row_i = icx // num_cells_image_height
11    cell_col_i = icx % num_cells_image_width
12
13    bbox_center_x_coord = cell_col_i * yolo_interval + yolo_interval/2 +
del_x * yolo_interval
14    bbox_center_y_coord = cell_row_i * yolo_interval + yolo_interval/2 +
del_y * yolo_interval
15    x1 = int(bbox_center_x_coord - w / 2.0)
16    y1 = int(bbox_center_y_coord - h / 2.0)
17
18    return [x1, y1, int(w + x1), int(h + y1)]
19
20 # Based on Professor Kak's run_code_for_testing_multi_instance_detection
function
21 def test(net, valLoader):
22     net = net.to(device)
23
24     with torch.no_grad():
25         for i, data in enumerate(valLoader):
26             batch_images, batch_bbox, batch_labels, batch_yolo_tensor = data
27
28             batch_images = batch_images.to(device)
29             batch_bbox = batch_bbox.to(device)
30             batch_labels = batch_labels.to(device)
31             batch_yolo_tensor = batch_yolo_tensor.to(device)
32
33             predicted_yolo_tensor = net(batch_images) # Get Prediction
34             predicted_yolo_tensor = predicted_yolo_tensor.view(batch_size,
num_yolo_cells, num_anchor_boxes, yolo_vector_size+1) # Save YOLO info into

```



```

35 tensor
36         for ibx in range(predicted_yolo_tensor.shape[0]): # Across batch
37 axis
38             icx_to_best_anchor_box = {ic: None for ic in range(
39 num_cells_image_height * num_cells_image_width)}
40             for icx in range(predicted_yolo_tensor.shape[1]): # Across yolo
41 cell axis
42                 cell_pred_i = predicted_yolo_tensor[ibx, icx]
43                 prev_best = 0
44
45                 #Compare
46                 for anch_i in range(cell_pred_i.shape[0]):
47                     if(cell_pred_i[anch_i][0] > cell_pred_i[prev_best][0]):
48                         prev_best = anch_i
49                 icx_to_best_anchor_box[icx] = prev_best
50
51                 # Get the 5 yolo cells
52                 sorted_icx_to_box = sorted(icx_to_best_anchor_box,
53 key=lambda x: predicted_yolo_tensor[ibx,x,
54 icx_to_best_anchor_box[x]][0].item(), reverse=True)
55                 retained_cells = sorted_icx_to_box[:5]
56
57                 # Identify the objects in the retained cells and extract their
58 bounding boxes
59                 predicted_bboxes = []
60                 predicted_labels = []
61
62                 for icx in retained_cells:
63                     predicted_yolo_vector = predicted_yolo_tensor[ibx, icx,
64 icx_to_best_anchor_box[icx]]
65                     target_yolo_vector = batch_yolo_tensor[ibx, icx,
66 icx_to_best_anchor_box[icx]]
67
68                     class_label_predictions = predicted_yolo_vector[-4:]
69                     class_labels_probs = torch.nn.Softmax(dim=0)(
70 class_label_predictions)
71                     class_labels_probs = class_labels_probs[:-1]
72                     if(torch.all(class_labels_probs < 0.2)):
73                         predicted_class_label = None
74
75                     else:
76                         # Get the predicted class label
77                         best_predicted_class_index = (class_labels_probs ==
78 class_labels_probs.max())
79                         best_predicted_class_index = torch.nonzero(
80 best_predicted_class_index, as_tuple=True)
81                         predicted_class_label = class_list[
82 best_predicted_class_index[0].item()]
83                         predicted_labels.append(predicted_class_label)
84
85                         # Analyze the predicted regression elements
86                         pred_regress_vec = predicted_yolo_vector[1:5].cpu()
87                         pred_bb = AnalyzePred(pred_regress_vec, icx)
88                         predicted_bboxes.append(pred_bb)
89
90

```

```

80         if(i % 50 == 49):
81             if(predicted_bboxes and predicted_labels):
82                 display_images(batch_images[ibx], batch_labels[ibx].
118 tolist(), batch_bbox[ibx],
83                                     predicted_labels, predicted_bboxes)

```

**Listing 6:** *Test*

```

1 # Get DataFrame Data
2 trainDataFrame = pd.read_csv("/scratch/gilbreth/dfarache/ece60146/David/HW6/
   train_labels.csv")
3 valDataFrame = pd.read_csv("/scratch/gilbreth/dfarache/ece60146/David/HW6/
   val_labels.csv")
4
5 # Datasets Objects
6 trainDataset = MyDataset(trainDataFrame, root=root)
7 valDataset = MyDataset(valDataFrame, root=root)
8
9 # Data Loaders
10 batch_size = 64
11 trainLoader = torch.utils.data.DataLoader(valDataset, batch_size=batch_size,
   num_workers=2, shuffle=True, drop_last=True)
12 valLoader = torch.utils.data.DataLoader(valDataset, batch_size=batch_size,
   num_workers=2, shuffle=True, drop_last=True)
13
14 # Create neural network
15 net = NetForYolo(skip_connections=True, depth=8)
16
17 # Train Network
18 BCE_loss, MSE_loss, CE_loss = train(net, trainLoader, epochs=8, lr=1e-5, betas
   =(0.9, 0.999))
19
20 plot_losses(BCE_loss, MSE_loss, CE_loss)
21
22 # Test Network
23 images = test(net, valLoader)

```

**Listing 7:** *Main for running all prior code*

## 7 Task 5: Validation and Results

The following images are the results and validation for the model created in the prior code. It will be shown that all models met the criteria of high enough accuracy in terms of loss values.

The loss in Figure 2 indicates proper learning by the model, but now looking at Figure 3 it may be overfitting. Figure three has a tendency to guess for buses and place 5 boundary boxes within a certain arrays. These boxes do move and adjust to the image, indicating that learning is happening but the results are not trust worthy and probably has some issue with the data-loader not properly iterating through. This would lead to overfitting, I had tried to find the error but failed to do so but the rest of the code works properly in its separate parts, just not its final results.




Figure 2: *Layer in network*




Figure 3: *Resulting images*

## 8 Lessons learned

From this assignment I learned YOLO method, struggling with the implementation but able to process and learn from it.