

HW 7: GAN

David E. Farache, Email ID: dfarache@purdue.edu

April 6, 2023

1 Introduction

The goal of this homework is to create a pizza-generative adversarial network (GAN). In this case, we experiment with both Deep Convolutions GAN (DCGAN) and a Wasserstein GAN (WGAN) which use different approximation methods to calculate the difference between images and reduce those to create proper fake images.

2 Explanation of Networks

2.1 GAN

An adversarial network refers to a method whereby a generator is utilized to create fake images that are then evaluated by a discriminator, with the goal being that the discriminator is no longer capable of distinguishing between real and fake images. In this case, we assume that each image is a probability distribution $\rho_x(x)$, where x is the image, and that each generated image, $G(z, \theta)$ where z is the noise vector and θ are the parameters, is a probability of $\rho_\theta(z)$. The goal is for the difference between distributions to be near zero, meaning that the probability of the image being true or false is about equivalent.

The generator is trained to minimize the ability of the discriminator to differentiate between the real and fake images:

$$\theta = \min_{\theta} E[\log(1 - D(G(z)))] \quad (1)$$

The discriminator is trained by maximizing the probability of giving the proper label to the image:

$$\theta = \max_{\theta} E[\log(D(x))] \quad (2)$$

This creates a minmax equation for the GAN network as seen below:

$$\min_{\theta_g} \max_{\theta_d} [E[\log(1 - D(G(z)))] + E[\log(D(x))]] \quad (3)$$

2.2 DCGAN

The deep convolution GAN, is a network that uses convolutional layers within both the generator and the discriminator. In the case the DCGAN is trained using Binary Cross-Entropy (BCE) loss with the prediction being a binary label of the real or fake image. This equation for the loss can be seen below:

$$BCE = -(y \log(p) + (1 - y) \log(1 - p)) \quad (4)$$

2.3 WGAN

Wasserstein GAN utilizes a different method from the previously mentioned networks, as it does not use a discriminator but a critic, which differs by the objective not to evaluate the generator results via a binary classification but instead calculate the Wasserstein distance between training and generated data distribution. Wasserstein distance is based on the marginal distribution. This is done via the equation below:

$$dw(P_r, P_\theta) = \sup_{\|f\|_L \leq 1} [E(f_w(x)) - E(f_w(g_\theta(z)))] \quad (5)$$

In this equation marginals of the real data to the generated data are compared for all 1-Lipschitz Functions described by:

$$|f(x_1) - f(x_2)| \leq 1 * d(x_1, x_2) \forall x_1, x_2 \in X \quad (6)$$

The designation of the WGAN is to reduce the Wasserstein distance between the distribution of the generated set and the real set. While doing so, critic C attempts to maximize the distance to increase the diversity of the images capable of being generated so as to not simply be a copier. This again leads to a min-max problem for the framework as described below:

$$\min_g \max_c [E[C(x)] - E[C(G(z))]] \quad (7)$$

This was further improved via the implementation of a gradient penalty, introducing a method to find optimal critic C by minimizing critic loss as seen below:

$$CriticLoss = E[C(G(z))] - E[C(x)] + [\|\nabla_{\hat{x}} C(\hat{x})\|^2 1]^2 \quad (8)$$

3 Task 1: DCGAN

3.1 Setup Code

```
1 # %%
2 # Libraries
3 import numpy as np
4 import torch
5 import torchvision.transforms as tvf
6 import torch.utils.data
7 import torch.nn as nn
8 import torch.nn.functional as F
9 import matplotlib.pyplot as plt
10 from PIL import Image
11 import os
12 from pprint import pprint
13 from torchinfo import summary
14 import torchvision.datasets
15 import time
16 import datetime
17 from pytorch_fid.fid_score import calculate_activation_statistics,
18     calculate_frechet_distance
19 from pytorch_fid.inception import InceptionV3
20 from torchvision.utils import save_image
```

```

20
21 device = 'cuda'
22 device = torch.device(device)
23 root_dir = "/scratch/gilbreth/dfarache/ece60146/David/HW7/"
24 train_data_path = root_dir + "pizza_train"
25 test_data_path = root_dir + "pizza_eval"
26
27 # Create Data Loader
28 def createDataLoader(root, batch_size, image_shape):
29     transform = tvl.Compose([tvl.Resize(image_shape),
30                             tvl.CenterCrop(image_shape),
31                             tvl.ToTensor(),
32                             tvl.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])
33     data_set = torchvision.datasets.ImageFolder(root, transform=transform)
34
35     DataLoader = torch.utils.data.DataLoader(data_set, batch_size=batch_size,
36 num_workers=2, shuffle=True, drop_last=True)
37     return DataLoader

```

Listing 1: Setup Code

3.2 DCGAN Network

```

1 # Discriminator
2 # Based on DLStudio Discriminator-Generator DG1
3 class Discriminator(nn.Module):
4     def __init__(self):
5         super(Discriminator, self).__init__()
6         # Conv Layers
7         self.conv_in = nn.Conv2d( 3, 64, kernel_size=4, stride=2, padding=1)
8         self.conv_in2 = nn.Conv2d( 64, 128, kernel_size=4, stride=2, padding=1)
9         self.conv_in3 = nn.Conv2d( 128, 256, kernel_size=4, stride=2, padding=1)
10        self.conv_in4 = nn.Conv2d( 256, 512, kernel_size=4, stride=2, padding=1)
11        self.conv_in5 = nn.Conv2d( 512, 1024, kernel_size=4, stride=2, padding
=1)
12        self.conv_in6 = nn.Conv2d( 1024, 1, kernel_size=4, stride=1, padding=1)
13
14        # Batch Layers
15        self.bn1 = nn.BatchNorm2d(128)
16        self.bn2 = nn.BatchNorm2d(256)
17        self.bn3 = nn.BatchNorm2d(512)
18        self.bn4 = nn.BatchNorm2d(1024)
19
20        # Sig
21        self.sig = nn.Sigmoid()
22
23        def forward(self, x):
24            x = torch.nn.functional.leaky_relu(self.conv_in(x), negative_slope=0.2,
inplace=True)
25            x = self.bn1(self.conv_in2(x))
26
27            x = torch.nn.functional.leaky_relu(x, negative_slope=0.2, inplace=True)
28            x = self.bn2(self.conv_in3(x))
29
30            x = torch.nn.functional.leaky_relu(x, negative_slope=0.2, inplace=True)
31            x = self.bn3(self.conv_in4(x))

```

```

32         x = torch.nn.functional.leaky_relu(x, negative_slope=0.2, inplace=True)
33         x = self.bn4(self.conv_in5(x))
34
35         x = torch.nn.functional.leaky_relu(x, negative_slope=0.2, inplace=True)
36         x = self.conv_in6(x)
37
38
39         x = self.sig(x)
40         return x
41
42 # Generator
43 # Based on slides in class
44 class Generator(nn.Module):
45     def __init__(self):
46         super(Generator, self).__init__()
47
48         # Conv Layers
49         self.latent_to_image = nn.ConvTranspose2d(100, 512, kernel_size=4,
50 stride=1, padding=0, bias=False)
51         self.upsampler2 = nn.ConvTranspose2d( 512, 256, kernel_size=4, stride=2,
52 padding=1, bias=False)
53         self.upsampler3 = nn.ConvTranspose2d( 256, 128, kernel_size=4, stride=2,
54 padding=1, bias=False)
55         self.upsampler4 = nn.ConvTranspose2d( 128, 64, kernel_size=4, stride=2,
56 padding=1, bias=False)
57         self.upsampler5 = nn.ConvTranspose2d( 64, 3, kernel_size=4, stride=2,
58 padding=1, bias=False)
59
60         # Batch Layers
61         self.bn1 = nn.BatchNorm2d(512)
62         self.bn2 = nn.BatchNorm2d(256)
63         self.bn3 = nn.BatchNorm2d(128)
64         self.bn4 = nn.BatchNorm2d(64)
65
66         # Tanh
67         self.tanh = nn.Tanh()
68
69     def forward(self, x):
70         x = self.latent_to_image(x)
71
72         x = torch.nn.functional.relu(self.bn1(x))
73         x = self.upsampler2(x)
74
75         x = torch.nn.functional.relu(self.bn2(x))
76         x = self.upsampler3(x)
77
78         x = torch.nn.functional.relu(self.bn3(x))
79         x = self.upsampler4(x)
80
81         x = torch.nn.functional.relu(self.bn4(x))
82         x = self.upsampler5(x)
83
84         x = self.tanh(x)
85         return x

```

Listing 2: DCGAN generator and discriminator

3.3 DCGAN Training

```
1 # Training DCGAN
2 def weights_init(m):
3     """
4     From the DCGAN paper, the authors specify that all model weights shall be
5     randomly initialized from a Normal distribution with mean=0, stdev=0.02.
6     The weights_init function takes an initialized model as input and
7     reinitializes
8     all convolutional, convolutional-transpose, and batch normalization layers
9     to
10    meet this criteria. This function is applied to the models immediately after
11    initialization.
12
13    https://pytorch.org/tutorials/beginner/dcgan\_faces\_tutorial.html
14    """
15
16    classname = m.__class__.__name__
17    if(classname.find('Conv') != -1): # If Conv not found in the classname
18        nn.init.normal_(m.weight.data, mean=0.0, std=0.02)
19    elif(classname.find('BatchNorm') != -1): # If BatchNorm not found in the
20        classname
21        nn.init.normal_(m.weight.data, mean=1.0, std=0.02)
22        nn.init.constant_(m.bias.data, val=0)
23
24 # Based on lecture slides
25 def TrainDCGAN(netD, netG, epochs, betas, lr, trainDataLoader):
26     # Number of channcel for noise vector
27     nz = 100
28
29     # Optimizer to device
30     netD = netD.to(device)
31     netG = netG.to(device)
32
33     # Apply weight
34     netD.apply(weights_init)
35     netG.apply(weights_init)
36
37     # We will use the same noise batch to periodically check on the progress
38     # made for the Generator:
39     fixed_noise = torch.randn(batch_size, nz, 1, 1, device=device)
40
41     # Establish convention for real and fake labels during training
42     real_label = 1
43     fake_label = 0
44
45     # Adam optimizers for the Discriminator and the Generator:
46     optimizerD = torch.optim.Adam(netD.parameters(), lr=lr, betas=betas) # Adam
47     # Optimizer for Discriminator
48     optimizerG = torch.optim.Adam(netG.parameters(), lr=lr, betas=betas) # Adam
49     # Optimizer for Generator
50
51     # Criterion BCE
52     criterion = nn.BCELoss()
53
54     # Lists for training data
55     img_list = []
```

```

50 G_losses = []
51 D_losses = []
52 iters = 0
53
54 print("\n\nStarting Training Loop...\n\n")
55 start_time = time.perf_counter()
56
57 for epoch in range(epochs):
58     g_losses_per_print_cycle = []
59     d_losses_per_print_cycle = []
60
61     for i, data in enumerate(trainDataLoader, 0):
62
63         # Get Real Images
64         netD.zero_grad()
65         real_images_in_batch = data[0].to(device)
66
67         # Train Discriminator on real images
68         label = torch.full((real_images_in_batch.size(0),), real_label,
69 dtype=torch.float, device=device)
70         output = netD(real_images_in_batch).view(-1)
71
72         real_image_Dlosses = criterion(output, label)
73         real_image_Dlosses.backward()
74
75         # Train Discriminator on fakes
76         noise = torch.randn(real_images_in_batch.size(0), nz, 1, 1, device=
77 device)
78         fakes = netG(noise) # Create fakes
79         label.fill_(fake_label) # Fill label with fakes
80
81         output = netD(fakes.detach()).view(-1) # Get outputs of
82 discriminator
83
84         fake_image_Dlosses = criterion(output, label)
85         fake_image_Dlosses.backward()
86         total_Dlosses = real_image_Dlosses + fake_image_Dlosses
87         d_losses_per_print_cycle.append(total_Dlosses)
88
89         optimizerD.step() # Only the Discriminator weights are incremented
90
91         # Minimize 1 - D(G(z)) by maximize D(G(z)) with generator of target
92 value 1
93         netG.zero_grad()
94
95         label.fill_(real_label)
96         output = netD(fakes).view(-1)
97
98         total_Glosses = criterion(output, label)
99         g_losses_per_print_cycle.append(total_Glosses)
100
101         total_Glosses.backward()
102         optimizerG.step()
103
104         if i % 100 == 99:
105             os.makedirs(root_dir + "./model", exist_ok = True)

```

```

103         mean_D_loss = torch.mean(torch.FloatTensor(
104             d_losses_per_print_cycle))
105         mean_G_loss = torch.mean(torch.FloatTensor(
106             g_losses_per_print_cycle))
107
108         print("[epoch=%d/%d iter=%4d elapsed_time=%5d secs] mean_D_loss
109             =%7.4f mean_G_loss=%7.4f" %
110             ((epoch+1), epochs, (i+1), time.time(), mean_D_loss, mean_G_loss))
111
112         d_losses_per_print_cycle = []
113         g_losses_per_print_cycle = []
114
115         torch.save(netG.state_dict(), root_dir + "./model/DCGAN_gen.pt")
116         torch.save(netD.state_dict(), root_dir + "./model/DCGAN_disc.pt"
117     )
118
119     with torch.no_grad():
120         fake = netG(fixed_noise).detach().cpu()
121         img_list.append(torchvision.utils.make_grid(fake, padding=1,
122             pad_value=1, nrow=4, normalize=True))
123
124     # Get All Losses
125     G_losses.append(total_Glosses.item())
126     D_losses.append(total_Dlosses.item())
127
128     print("Traing Time %s sec" % (time.time() - start_time))
129     return G_losses, D_losses, img_list

```

Listing 3: DCGAN training

4 Task 2: WGAN

4.1 WGAN Network

For WGAN the same generator was used as that in DCGAN

```

1 # WGAN
2 # Based on DLStudio Critic-Generator CG2
3 class Critic(nn.Module):
4     def __init__(self):
5         super(Critic, self).__init__()
6         self.DIM = 64
7         self.net = nn.Sequential(
8             nn.Conv2d(3, self.DIM, kernel_size=5, stride=2, padding=2),
9             nn.ReLU(True),
10            nn.Conv2d(self.DIM, 2*self.DIM, kernel_size=5, stride=2, padding=2),
11            nn.ReLU(True),
12            nn.Conv2d(2*self.DIM, 4*self.DIM, kernel_size=5, stride=2, padding
13            =2),
14            nn.ReLU(True),
15            nn.Conv2d(4*self.DIM, 8*self.DIM, kernel_size=5, stride=2, padding
16            =2),
17            nn.ReLU(True),
18        )
19        self.output = nn.Linear(4*4*4*self.DIM, 1)

```

```

19     def forward(self, x):
20         x = x.view(-1, 3, 64, 64)
21         x = self.net(x)
22
23         x = x.view(-1, 4*4*4*self.DIM)
24         x = self.output(x)
25
26         x = x.mean(0)
27         x = x.view(1)
28         return x

```

Listing 4: WGAN generator and critic

4.2 WGAN Training

```

1  # WGAN
2  # Based on lecture slides
3  def calc_gradient_penalty(netC, real_data, fake_data, LAMBDA=10):
4      epsilon = torch.rand(1).cuda()
5
6      interpolates = epsilon * real_data + ((1 - epsilon) * fake_data)
7      interpolates = interpolates.requires_grad_(True).cuda()
8
9      critic_interpolates = netC(interpolates)
10
11     gradients = torch.autograd.grad(outputs = critic_interpolates, inputs=
interpolates,
12                                     grad_outputs = torch.ones(critic_interpolates.size()
).cuda(),
13                                     create_graph = True, retain_graph=True,
14                                     only_inputs=True)[0]
15
16     gradient_penalty = ((gradients.norm(2, dim=1) - 1) ** 2).mean() * LAMBDA
17     return gradient_penalty
18
19 # Based on lecture slides
20 def TrainWGAN(netC, netG, epochs, betas, lr, trainloader):
21     nz = 100 # Set the number of channels for the 1x1 input noise vectors for
the Generator
22
23     netG = netG.to(device)
24     netC = netC.to(device)
25
26     netG.apply(weights_init) # initialize network parameters
27     netC.apply(weights_init) # initialize network parameters
28
29     fixed_noise = torch.randn(batch_size, nz, 1, 1, device=device) # Make noise
vector
30
31     one = torch.tensor([1], dtype=torch.float).to(device)
32     minus_one = torch.tensor([-1], dtype=torch.float).to(device)
33
34     # Adam optimizers
35     optimizerC = torch.optim.Adam(netC.parameters(), lr=lr, betas=betas)
36     optimizerG = torch.optim.Adam(netG.parameters(), lr=lr, betas=betas)
37

```



```

38 img_list = []
39 G_losses = []
40 C_losses = []
41
42 iters = 0
43 gen_iterations = 0
44
45 print(f"Training started at time {datetime.datetime.now().time()}")
46 start_time = time.time()
47
48 for epoch in range(epochs):
49
50     data_iter = iter(trainDataLoader)
51     i = 0
52     ncritic = 5
53
54     while i < len(trainDataLoader):
55
56         for p in netC.parameters():
57             p.requires_grad = True
58             ic = 0
59
60         while ic < ncritic and i < len(trainDataLoader):
61             ic += 1
62
63             # Training with real images
64             netC.zero_grad()
65
66             real_images_in_batch = next(data_iter)
67             real_images_in_batch = real_images_in_batch[0].to(device)
68
69             i += 1
70
71             # Mean value for all images
72             critic_for_reals_mean = netC(real_images_in_batch)
73
74             # Target gradient -1
75             critic_for_reals_mean.backward(minus_one)
76
77             # Train with fake images
78             noise = torch.randn(real_images_in_batch.size(0), nz, 1, 1,
device=device)
79             fakes = netG(noise)
80
81             # Mean value for batch
82             critic_for_fakes_mean = netC(fakes.detach())
83
84             # Aim for target of 1
85             critic_for_fakes_mean.backward(one)
86
87             #Gradient penalty
88             gradient_penalty = calc_gradient_penalty(netC,
real_images_in_batch, fakes)
89             gradient_penalty.backward()
90
91             # Calc distance
92             wasser_dist = critic_for_reals_mean - critic_for_fakes_mean

```

```

93         loss_critic = -wasser_dist + gradient_penalty
94
95         # Update the Critic
96         optimizerC.step()
97
98         for p in netC.parameters():
99             p.requires_grad = False
100
101         # Train generator
102         netG.zero_grad()
103
104         noise = torch.randn(real_images_in_batch.size(0), nz, 1, 1, device=
device)
105         fakes = netG(noise)
106
107         critic_for_fakes_mean = netC(fakes)
108         loss_gen = critic_for_fakes_mean
109         critic_for_fakes_mean.backward(minus_one)
110
111         # Update the Generator
112         optimizerG.step()
113         gen_iterations += 1
114
115         if i % (ncritic * 20) == 0:
116             os.makedirs(root_dir + "./models", exist_ok = True)
117
118             print("[epoch=%d/%d iter=%4d elapsed_time=%5d secs] mean_C_loss
=%7.4f mean_G_loss=%7.4f wass_dist=%7.4f" %
119                 ((epoch+1), epochs, (i+1), time.time(), loss_critic.data[0],
loss_gen.data[0], wasser_dist.data[0]))
120
121             torch.save(netG.state_dict(), root_dir + "./model/WGAN_gen.pt")
122             torch.save(netC.state_dict(), root_dir + "./model/WGAN_crit.pt")
123
124         # Get All Losses
125         G_losses.append(loss_gen.data[0].item())
126         C_losses.append(loss_critic.data[0].item())
127
128         with torch.no_grad():
129             fake_image = netG(fixed_noise).detach().cpu()
130             img_list.append(torchvision.utils.make_grid(fake_image, padding=1,
pad_value=1, nrow=4, normalize=True))
131
132
133     print("Traing Time %s sec" % (time.time() - start_time))
134     return G_losses, C_losses, img_list

```

Listing 5: *WGAN training*

5 Task 3: Evaluation

```

1 # Plotting
2 def plotLoss(lossGen, lossDis, epochs):
3     # Plot the training losses
4     iterations = range(len(lossDis))

```

```

5
6 fig = plt.figure(1)
7 plt.plot(iterations, lossGen, label="Generator Loss")
8 plt.plot(iterations, lossDis, label="Discriminator Loss")
9
10 plt.legend()
11
12 plt.xlabel("Iterations", fontsize = 16)
13 plt.ylabel("Loss", fontsize = 16)
14
15 plt.show()

```

Listing 6: *Plotting Results*

```

1 # Save Fake Images
2
3 def ProduceFakes(netG, netGW):
4     # Make dir if not there
5     os.makedirs(root_dir + "/DCGAN_fakes", exist_ok = True)
6     os.makedirs(root_dir + "/WGAN_fakes", exist_ok = True)
7
8     #Noise
9     num_images = 1000
10
11     # Load Networks
12     netG.load_state_dict(torch.load(os.path.join(root_dir, "model/DCGAN_gen.pt")
13 ))
14     netG.eval()
15
16     netGW.load_state_dict(torch.load(os.path.join(root_dir, "model/WGAN_gen.pt")
17 ))
18     netGW.eval()
19
20     #fake_images = generated_fake_images.detach().cpu()
21
22     #for i in range(num_images):
23         validation_noise = torch.randn(num_images, 100, 1, 1, device=device)
24
25         # Generate image with dcgan and wgan
26         dcgan_img = netG(validation_noise)
27         wgan_img = netGW(validation_noise)
28
29         # Convert to cpu
30         fake_dcgan_images = dcgan_img.detach().cpu()
31         fake_wgan_images = wgan_img.detach().cpu()
32
33         for i in range(len(fake_dcgan_images)):
34             image_dcgan = tv.t.ToPILImage()(fake_dcgan_images[i] / 2 + 0.5)
35             image_dcgan.save(os.path.join(root_dir + "./DCGAN_fakes/", "DCGAN_image_
36 {0}.png".format(i+1)))
37
38             image_wgan = tv.t.ToPILImage()(fake_wgan_images[i] / 2 + 0.5)
39             image_wgan.save(os.path.join(root_dir + "./WGAN_fakes/", "WGAN_image_
40 {0}.png".format(i+1)))

```

Listing 7: *Producing Fake Images for FID*

```

1 # Validation
2 def compute_fid_score(fake_paths, real_paths, dims=2048):
3     block_idx = InceptionV3.BLOCK_INDEX_BY_DIM[dims]
4     model = InceptionV3([block_idx]).to(device)
5
6     m1, s1 = calculate_activation_statistics( real_paths, model, device=device)
7     m2, s2 = calculate_activation_statistics( fake_paths, model, device=device)
8
9     fid_value = calculate_frechet_distance(m1, s1, m2, s2)
10    print(f'FID: {fid_value:.2f}')
11
12    return fid_value

```

Listing 8: *FID Score Calculator*

6 Task 4: Results and Plots

Plotting for loss over iteration and the resulting image predictions are shown below.

```

1 # Main
2
3 # Get Trainloader
4 batch_size = 16
5 image_shape = 64
6
7 trainDataLoader = createDataLoader(train_data_path, batch_size, image_shape)
8 testDataLoader = createDataLoader(test_data_path, batch_size, image_shape)
9
10 images, _ = next(iter(trainDataLoader))
11 plt.figure(figsize=(10,10))
12 plt.axis("off")
13 plt.title("Training Images")
14 plt.imshow(np.transpose(torchvision.utils.make_grid(images[:batch_size],
15                                                     padding=2, normalize=True)
16                                                     ,(1,2,0)))
17
18 # Set Params
19
20 lr = 1e-5*4
21 betas = (0.5, 0.999)
22 epochs = 87
23
24 netG = Generator()
25 netD = Discriminator()
26
27 G_losses, D_losses, DCGAN_image_list = TrainDCGAN(netD, netG, epochs, betas, lr,
28                                                    trainDataLoader)
29
30 netGW = Generator()
31 netC = Critic()
32
33 lr = 1e-3
34 epochs = 250
35
36 GW_losses, C_losses, WGAN_image_list = TrainWGAN(netC, netGW, epochs, betas, lr,
37                                                    trainDataLoader)

```

```

35
36 plotLoss(G_losses, D_losses, epochs) #DCGAN
37
38 plotLoss(GW_losses, C_losses, epochs) #WGAN
39
40 plot_fake_real_images(DCGAN_image_list, trainDataLoader) # DCGAN
41
42 plot_fake_real_images(WGAN_image_list, trainDataLoader) # WGAN
43
44 ProduceFakes(netG, netGW)
45
46 real_imgs = [test_data_path + "/eval/" + i for i in os.listdir(test_data_path +
47     "/eval/")]
48 dcgan_imgs = [root_dir + "./DCGAN_fakes/" + i for i in os.listdir(root_dir + "./
49     DCGAN_fakes/")]
50 wgan_imgs = [root_dir + "/WGAN_fakes/" + i for i in os.listdir(root_dir + "./
51     WGAN_fakes/")]
52
53 dcgan_fid = compute_fid_score(dcgan_imgs, real_imgs, dims=2048)
54 wgan_fid = compute_fid_score(wgan_imgs, real_imgs, dims=2048)

```

Listing 9: *Main To Run Code*

The results indicate that the WCGAN performs better than the DCGAN network. This is evident by the FID score being lower than $110 < 150$ and the images being more diverse and clear. This result makes sense as the WCGAN uses a more sophisticated method of the Wasserstein distance paired with the critic that is just a discriminator. It can be seen in Figure 3 vs Figure 2 that the images are less blurry and have fewer replication effects than the DCGAN. Furthermore, as seen in Figure 1, the loss of WCGAN is much lower and with less noise than that of DCGAN which spikes, significantly more which could be due to the issue of the vanishing gradient problem that more may have been resolvable by adding skip connections.




Figure 1: *WCGAN*

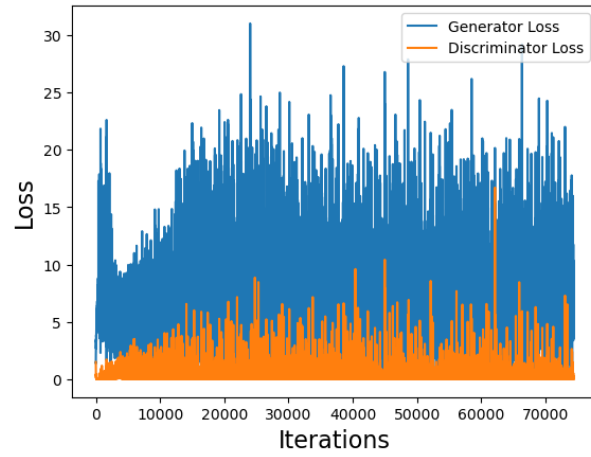


Figure 2: *DCGAN*

Figure 3: *Loss per iterations*



Figure 4: *DCGAN Fake and Real Images*




Figure 5: *WGAN Fake and Real Images*

```

1 dcgan_fid = compute_fid_score(dcgan_imgs, real_imgs, dims=2048)
100%|██████████| 20/20 [00:04<00:00, 4.70it/s]
100%|██████████| 20/20 [00:01<00:00, 10.10it/s]

FID: 150

1 wgan_fid = compute_fid_score(wgan_imgs, real_imgs, dims=2048)
100%|██████████| 20/20 [00:01<00:00, 10.31it/s]
100%|██████████| 20/20 [00:02<00:00, 9.91it/s]

FID: 110.46

```

Figure 6: *FID Score*

7 Lessons learned

From this assignment, I learned how to create a DCGAN and WGAN network along with creating a proper discriminator and critic for producing proper results.